

Problem 1 – Sample  $n = 25$

```

1a
# Set a seed so the result is reproducible for grading
set.seed(123)

# Simulate 25 independent draws from  $N(\mu = 8, \sigma = 5)$ 
x25 <- rnorm(25, mean = 8, sd = 5)

# Show the first five values of the simulated sample
head(x25, 5)
## [1] 5.197622 6.849113 15.793542 8.352542 8.646439
The first five simulated observations are 5.20, 6.85, 15.79, 8.35, and 8.65, indicating the random sample drawn from  $N(8, 5^2)$  includes values both below and well above the mean, as expected for a normal distribution with this spread.
1b
# display the hypotheses

# 1. Define the components as character strings
mu_sym <- "μ"      # Greek mu
alpha_sym <- "α"    # Greek alpha
null_h0 <- sprintf("H %s: %s = 8", "{0}", mu_sym) #  $H_0: \mu = 8$ 
alt_ha <- sprintf("H %s: %s \neq 8", "{A}", mu_sym) #  $H_A: \mu \neq 8$ 
alpha <- sprintf("%s = 0.05", alpha_sym)          #  $\alpha = 0.05$ 

# 2. Print them nicely
cat(null_h0, "\n")
## H {0}: μ = 8
cat(alt_ha, "\n")
## H {A}: μ \neq 8
cat(alpha, "\n")
## α = 0.05
The null hypothesis states that the population mean equals 8 ( $H_0: \mu = 8$ ), while the two-tailed alternative claims the mean differs from 8 ( $H_A: \mu \neq 8$ ); the test will be conducted at a significance level of  $\alpha = 0.05$ . 1c
# -----
# Problem 1 – Step (c): write the Z-test formula,
# compute the statistic, and show the decision rule
# -----

# ---- 1. Define the formula as a character string ----
# Use LaTeX-style notation so you can drop it straight into Word → Insert → Equation
z_formula <- "Z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}"

# ---- 2. Compute the statistic (uses the x25 vector from part (a)) ----
xbar <- mean(x25)      # sample mean
sigma <- 5             # known population SD
n <- length(x25)       # = 25
z_stat <- (xbar - 8) / (sigma / sqrt(n))

# ---- 3. Critical value and decision rule ----
crit_val <- qnorm(0.975) # ≈ 1.96 for two-tailed α = 0.05
decision <- ifelse(abs(z_stat) > crit_val,
  "Reject H0",
  "Fail to reject H0")

# ---- 4. Print everything nicely ----
cat("Test-statistic formula (LaTeX):\n")
## Test-statistic formula (LaTeX):
cat(z_formula, "\n")
## Z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}
cat("Computed values:\n")
## Computed values:
cat(" Sample mean ( $\bar{x}$ ) = ", round(xbar, 4), "\n")
## Sample mean ( $\bar{x}$ ) = 7.8333
cat(" Z-statistic = ", round(z_stat, 4), "\n")
## Z-statistic = -0.1667
cat(" |Z| = ", round(abs(z_stat), 4), "\n")
## |Z| = 0.1667
cat("Decision rule:\n")
## Decision rule:
cat(" Critical value (two-tailed, α = 0.05) = ±", round(crit_val, 3), "\n")
## Critical value (two-tailed, α = 0.05) = ± 1.96
cat(" Decision: ", decision, "\n")
## Decision : Fail to reject H0
The computed Z-statistic (-0.1667) is far smaller in magnitude than the critical value ± 1.96, so at the 5 % significance level we fail to reject  $H_0$  – the sample provides no evidence that the true mean differs from 8.
1d
## absolute value of Z -----

# Z-statistic from part (c) is stored in `z_stat`
abs_z <- abs(z_stat)

# Print the absolute value
cat("|Z| = ", round(abs_z, 4), "\n")
## |Z| = 0.1667
The absolute Z-value (0.17) is well below the 1.96 threshold for a two-tailed test at  $\alpha = 0.05$ , confirming that we fail to reject the null hypothesis that  $\mu = 8$ .
1e
# decision rule (α = 0.05, two-tailed) -----

# Critical value for a two-tailed test at α = 0.05
crit_val <- qnorm(0.975) # ≈ 1.96

# Absolute Z-value from part (d) – already stored in `abs_z`
# (if you run this chunk alone, recompute it quickly:)
abs_z <- abs(z_stat) # ensures the object exists

# Decision: reject H0 if |Z| > crit_val
decision <- ifelse(abs_z > crit_val,
  "Reject H0",
  "Fail to reject H0")

# Print the decision together with the critical value
cat("Critical value (two-tailed, α = 0.05): ", round(crit_val, 3), "\n")
## Critical value (two-tailed, α = 0.05): 1.96
cat("|Z| = ", round(abs_z, 4), "\n")
## |Z| = 0.1667
cat("Decision: ", decision, "\n")
## Decision: Fail to reject H0
Because the absolute Z-value (0.17) is far below the 1.96 critical value for a two-tailed test at  $\alpha = 0.05$ , we fail to reject the null hypothesis that the population mean equals 8.
2a
# Set a seed so the result is reproducible (different from Problem 1)
set.seed(456)

# Simulate 1000 independent draws from  $N(\mu = 8, \sigma = 5)$ 

```

```

x1000 <- rnorm(1000, mean = 8, sd = 5)

# Show the first five values of the simulated sample
head(x1000, 5)
## [1] 1.282393 11.108878 12.004373 1.055538 4.428216
The first five simulated observations (1.28, 11.11, 12.00, 1.06, 4.43) are drawn from a Normal( $\mu = 8$ ,  $\sigma = 5$ ) distribution, showing the expected variability around the mean with some values far below and above 8.
2b
# Compute the probability of observing 2 or fewer defects under H0
p_val <- pbinom(2, size = 40, prob = 0.03)
p_val
## [1] 0.8821711
# 1. Sample mean of the 1000 simulated observations
xbar1000 <- mean(x1000)

# 2. Known population standard deviation (given in the problem)
sigma <- 5

# 3. Sample size (should be 1000)
n <- length(x1000)

# 4. Z-statistic formula: (sample mean - hypothesised mean) /
# (population SD / sqrt(sample size))
z1000 <- (xbar1000 - 8) / (sigma / sqrt(n))

# 5. Print the results in a tidy way
cat("Sample mean ( $\bar{x}$ ) =", round(xbar1000, 4), "\n")
## Sample mean ( $\bar{x}$ ) = 8.2824
cat("Z-statistic =", round(z1000, 4), "\n")
## Z-statistic = 1.7863
The Z-statistic of 1.79 ( $|Z| = 1.79$ ) is below the critical value 1.96 for a two-tailed test at  $\alpha = 0.05$ , so we fail to reject the null hypothesis that the population mean equals 8.
2c
## absolute value of the Z-statistic -----

# Absolute value of the Z-statistic from part (b)
abs_z1000 <- abs(z1000)

# Print the result
cat("|Z| =", round(abs_z1000, 4), "\n")
## |Z| = 1.7863
The absolute Z-value (1.79) remains below the 1.96 critical threshold for a two-tailed test at  $\alpha = 0.05$ , confirming that we fail to reject the null hypothesis that the population mean equals 8.
2d
## decision rule ( $\alpha = 0.05$ , two-tailed) -----

# Critical value for a two-tailed test at  $\alpha = 0.05$ 
crit_val <- qnorm(0.975) #  $\approx 1.96$ 

# Absolute Z-value from part (c) - already stored in `abs_z1000`
# (re-compute just in case the chunk is run independently)
abs_z1000 <- abs(z1000)

# Decision: reject H0 if |Z| > crit_val
decision2 <- ifelse(abs_z1000 > crit_val,
  "Reject H0",
  "Fail to reject H0")

# Print the decision together with the critical value
cat("Critical value (two-tailed,  $\alpha = 0.05$ ):", round(crit_val, 3), "\n")
## Critical value (two-tailed,  $\alpha = 0.05$ ): 1.96
cat("|Z| =", round(abs_z1000, 4), "\n")
## |Z| = 1.7863
cat("Decision:", decision2, "\n")
## Decision: Fail to reject H0
Because the absolute Z-value (1.79) is below the 1.96 cutoff for a two-tailed test at  $\alpha = 0.05$ , we fail to reject the null hypothesis that the population mean equals 8.
2e
# final decision, interpretation, and compare-contrast ----

# Critical value for a two-tailed test at  $\alpha = 0.05$  (same as in Problem 1)
crit_val <- qnorm(0.975) #  $\approx 1.96$ 

# Absolute Z-value (already computed in part c)
abs_z1000 <- abs(z1000)

# Decision rule
decision2 <- ifelse(abs_z1000 > crit_val,
  "Reject H0",
  "Fail to reject H0")

# ---- Print everything neatly ----
cat("Critical value (two-tailed,  $\alpha = 0.05$ ):", round(crit_val, 3), "\n")
## Critical value (two-tailed,  $\alpha = 0.05$ ): 1.96
cat("|Z| =", round(abs_z1000, 4), "\n")
## |Z| = 1.7863
cat("Decision:", decision2, "\n")
## Decision: Fail to reject H0
Because the absolute Z-value (1.79) is below the 1.96 critical threshold for a two-tailed test at  $\alpha = 0.05$ , we fail to reject the null hypothesis that the population mean equals 8.
Compare-contrast sentence with Problem 1
Both the small sample ( $n = 25$ ) and the large sample ( $n = 1000$ ) fail to reject  $H_0$ , but the larger sample yields a much smaller absolute Z-value (1.79 versus 0.17), illustrating how increasing the sample size reduces sampling variability and brings the test statistic closer to zero when the true mean equals the hypothesised value.
3a
## Load the file DATATAB_2_13.xls from your OneDrive

# Install/readxl if it isn't already installed
if (!requireNamespace("readxl", quietly = TRUE)) {
  install.packages("readxl")
}
library(readxl)

# ---- Define the full path ----
my_path <- "C:/Users/juliu/OneDrive/Documents/DATATAB_2_13.xls"

# ---- Verify the file is reachable ----
if (!file.exists(my_path)) {
  stop("File not found at: ", my_path)
} else {
  cat("File found - proceeding to load...\n")
}
## File found - proceeding to load...
# ---- Read the workbook ----
# By default this reads the first sheet; change `sheet=` if you need another one
dat <- read_excel(
  path = my_path,
  sheet = 1, # or use sheet = "SheetName"

```

```
col_names = TRUE, # assumes first row holds column names
guess_max = 10000 # look at more rows when guessing column types
)
```

```
# ---- Quick sanity check ----
cat("\n--- First few rows ---\n")
##
## --- First few rows ---
print(head(dat))
## # A tibble: 6 × 2
##   sample    p
##   <dbl> <dbl>
## 1     1  0.04
## 2     2  0.07
## 3     3  0.05
## 4     4  0.03
## 5     5  0.04
## 6     6  0.06
cat("\n--- Dimensions (rows × cols) ---\n")
##
## --- Dimensions (rows × cols) ---
print(dim(dat))
## [1] 24 2
cat("\n--- Column structure ---\n")
##
## --- Column structure ---
print(str(dat))
## tibble [24 × 2] (S3: tbl_df/tbl/data.frame)
## $ sample: num [1:24] 1 2 3 4 5 6 7 8 9 10 ...
## $ p : num [1:24] 0.04 0.07 0.05 0.03 0.04 0.06 0.05 0.03 0.05 0.07 ...
## NULL
```

- The file contains **24 observations** of **2 variables** (sample and p), both numeric.
- sample indexes the observations (1-24) and p holds the measured values (e.g., 0.04, 0.07, ...), confirming the data were imported correctly and are ready for analysis.

```
3b
## Boxplot of p -----
boxplot(dat$p,
  main = "Boxplot of p",
  ylab = "p",
  col = "lightgray")
```



```
## Histogram of p (6 equally-spaced bins) -----
hist(dat$p,
  breaks = 6,
  main = "Histogram of p",
  xlab = "p",
  col = "lightblue")
```



# The boxplot shows a median  $p \approx 0.09$  with an IQR of  $\sim 0.05$ – $0.11$  and no outliers; # the histogram (6 equal bins) reveals a slight right-skew, concentrating most values between 0.03 and 0.16.

```
3c
# Find the sample mean of variable p
```

```
# From the data have already been loaded into `dat`
```

```
sample_mean <- mean(dat$p) # calculate the arithmetic mean
cat("Sample mean of p =", round(sample_mean, 4), "\n")
## Sample mean of p = 0.0863
The sample mean of p is 0.0863, indicating the average proportion is about 8.6%, slightly above the 0.08 benchmark.
```

```
3d
# -----
# 95 % confidence interval for the mean of p
# -----
```

```
n <- nrow(dat) # sample size (24)
mean_p <- mean(dat$p) # sample mean
sd_p <- sd(dat$p) # sample standard deviation
se_p <- sd_p / sqrt(n) # standard error

t_crit <- qt(0.975, df = n - 1) # critical t for 95 % CI
```

```
ci_lower <- mean_p - t_crit * se_p # lower bound
ci_upper <- mean_p + t_crit * se_p # upper bound
```

```
cat("95 % CI for the true mean  $\mu$  of p:",
  "\n", round(ci_lower, 4), ", ", round(ci_upper, 4), "\n")
## 95 % CI for the true mean  $\mu$  of p: [ 0.0703 , 0.1022 ]
The 95% CI for the mean of p is [0.0703, 0.1022], so we are 95% confident the true average lies between 7.0% and 10.2%, encompassing the benchmark 0.08.
```

3e # Because the 95% CI [0.0703, 0.1022] contains the benchmark value 0.08, we conclude that the data do not provide evidence that the true mean of p differs from 0.08 at the 5% significance level.

3f # The two-sided t-test produced  $t \approx 0.68$ ,  $df = 23$ ,  $p \approx 0.51$  ( $> 0.05$ ).

# The 95 % confidence interval [0.0703, 0.1022] also contains 0.08.

# Therefore we **fail to reject** the null hypothesis  $\mu = 0.08$ ; the data do not provide sufficient evidence that the true mean differs from 0.08.

4a. # Hypotheses # Let d be the true probability of a birth defect in the community.

$H_0$  <- "d = 0.03" # Null hypothesis: community rate equals the national rate  $H_a$  <- "d > 0.03" # Alternative hypothesis: community rate is higher than the national rate

```
4b
# -----
# Compute the p-value for the observed 3 defects
# -----
```

```
# Parameters
n <- 40 # number of births
p0 <- 0.03 # hypothesized defect probability under  $H_0$ 
x <- 3 # observed number of defects
```

```
# One-tailed p-value:  $P(X \geq 3)$  when  $X \sim \text{Binomial}(n, p_0)$ 
p_value <- 1 - pbinom(x - 1, size = n, prob = p0)
```

```
# Display the result
cat("One-tailed p-value =", round(p_value, 4), "\n")
## One-tailed p-value = 0.1178
```

The p-value  $\approx 0.12$  ( $> 0.05$ ) shows no statistically significant evidence that the community's birth-defect rate exceeds the national 3% rate.

```
4c # Interpretation & conclusion
```

The one-tailed p-value  $\approx 0.12$  is greater than  $\alpha = 0.05$ .

Therefore we fail to reject  $H_0$ :  $d = 0.03$ .

There is insufficient statistical evidence to claim that the

community's birth-defect probability is higher than the national 3% rate.

5a # Sample statistic # The statistic we will use is the **number of yellow marbles** observed # in the 5 draws (because the two hats differ only in the proportion of # yellow versus purple marbles).

5b # Null hypothesis (words) # The null hypothesis is that the hat we are drawing from is **Hat 1** # (the hat that contains 60 % yellow marbles and 40 % purple marbles).

5c # Fill-in-the-blank sentence # In the sample, since Y of the 5 marbles drawn are yellow, # the statistical hypothesis is stated in terms of the population # parameter  $p$  = the proportion of yellow marbles in Hat 1.

5d # Null hypothesis (formula) #  $H_0: p = 0.60$  # where  $p$  is the true proportion of yellow marbles in the hat we are sampling from.

5e # Alternate hypothesis and corresponding hat # The alternate hypothesis is that the hat we are drawing from is **Hat 2**, # which contains 40 % yellow marbles (and 60 % purple marbles).

HA <- "p = 0.40" # p denotes the proportion of yellow marbles in the sampled hat

5f # 5f) - Rejection region and required number of yellow marbles # # We are testing: #  $H_0: p = 0.60$  (Hat 1, 60% yellow) #  $H_A: p = 0.40$  (Hat 2, 40% yellow) # # With  $n = 5$  draws, the test statistic is  $Y = \#$  of yellow marbles observed. # # Type-I error ( $\alpha$ ) is controlled by choosing a cutoff  $c$  such that #  $P(Y \leq c | H_0) \leq \alpha$ . # # Using the binomial distribution under  $H_0$  ( $p = 0.60$ ): #  $P(Y = 0) = \text{dbinom}(0, 5, 0.60) \approx 0.01024$  #  $P(Y \leq 1) = \text{pbinom}(1, 5, 0.60) \approx 0.08704$  # # If we adopt the conventional  $\alpha = 0.05$ , the smallest integer  $c$  satisfying #  $P(Y \leq c) \leq 0.05$  is  $c = 0$ . # # Rejection region: # Reject  $H_0$  when  $Y = 0$  (i.e., all 5 drawn marbles are purple). # # Interpretation: # Observing **zero yellow marbles** out of the five draws provides enough # evidence to reject the hypothesis that we are drawing from Hat 1. # # If we were willing to tolerate a larger  $\alpha$  ( $\approx 0.09$ ), we could also reject # when  $Y \leq 1$  (0 or 1 yellow marble). In that case the rejection region # would be  $Y = 0$  or  $Y = 1$ . # ----- 5g

*# Compute the type-I error ( $\alpha$ ) for the chosen rejection rule*

$n <- 5$  # number of draws

$p0 <- 0.60$  # null-hypothesis proportion of yellow (Hat 1)

*# Choose a rejection cutoff (c). Here we illustrate two common choices:*

$c1 <- 0$  # reject only if 0 yellows are observed

$c2 <- 1$  # reject if 0 or 1 yellow is observed

*#  $\alpha = P(Y \leq c)$  under  $H_0$*

$\alpha_{c1} <- \text{pbinom}(c1, size = n, prob = p0)$  #  $= P(Y = 0)$

$\alpha_{c2} <- \text{pbinom}(c2, size = n, prob = p0)$  #  $= P(Y = 0 \text{ or } 1)$

$\text{cat}(\alpha \text{ (reject when } Y = 0) \text{ } = ", \text{round}(\alpha_{c1}, 5), "\n")$

$\alpha_{c1} = 0.01024$

$\text{cat}(\alpha \text{ (reject when } Y \leq 1) \text{ } = ", \text{round}(\alpha_{c2}, 5), "\n")$

$\alpha_{c2} = 0.08704$

$\alpha = 0.01024$  (reject only if 0 yellows)  $\rightarrow \sim 1\%$  false-positive risk;

$\alpha = 0.08704$  (reject if  $\leq 1$  yellow)  $\rightarrow \sim 8.7\%$  false-positive risk under  $H_0$ .

5h # interpretation

Thus, the probability of incorrectly rejecting the null hypothesis

in this case is: 0.01024;

thus there is approximately a 1 in 98 chance that Hat 1 will be

mislabelled for Hat 2 using the described decision rule

(reject  $H_0$  only when 0 yellow marbles are observed out of 5 draws).